

## Question 1. Data Exploration


### a. Description and Data Cleaning


The subject of this report is a list of properties that are currently listed for sale in the state of New York.

	bed	bath	acre_lot	zip_code	house_size	price	is_prev_sold	log_price
count	84040.000000	83946.000000	84040.000000	84036.000000	84040.000000	8.404000e+04	388925.000000	84040.000000
mean	3.930200	2.980690	10.624121	10983.073409	2472.279938	1.274604e+06	0.145542	13.612734
std	2.062923	1.756449	849.058032	688.870753	2326.090138	2.312462e+06	0.352647	0.836738
min	1.000000	1.000000	0.000000	6390.000000	122.000000	2.000000e+04	0.000000	9.903488
25%	3.000000	2.000000	0.060000	10514.000000	1370.000000	5.280000e+05	0.000000	13.176852
50%	4.000000	3.000000	0.140000	10916.000000	2000.000000	7.545000e+05	0.000000	13.533811
75%	5.000000	4.000000	0.570000	11233.000000	2880.000000	1.220000e+06	0.000000	14.014361
max	42.000000	43.000000	100000.000000	14534.000000	112714.000000	1.690000e+08	1.000000	18.945409

Figure 1: Summary Statistics of DataFrame (df)

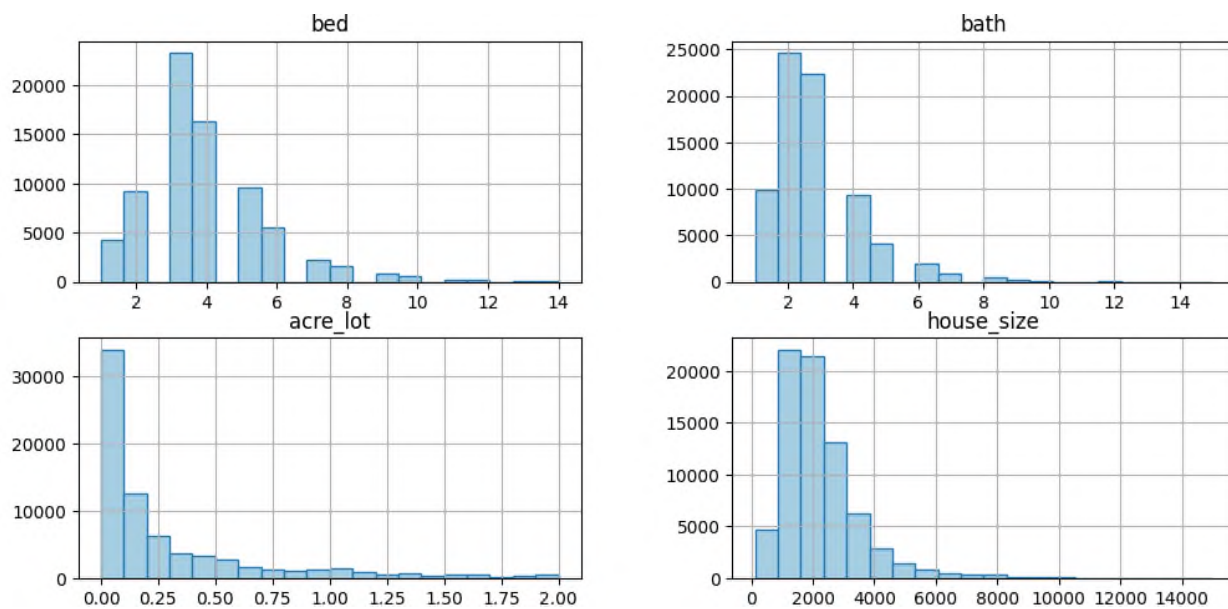
From the summary statistics in Figure 1, we found that 84,040 entries are present in the dataset. It includes information on the number of bedrooms and bathrooms, lot size, zip code, house size, sale price, whether the property was previously sold, and the natural logarithm of the sale price. Some missing values were observed under numerous categories, such as "bath", the number of bathrooms in a property, and "prev\_sold\_date", the last recorded date of property sale. The number of bedrooms ranged from 1 to an outlier of 42 with an average of 3-4 bedrooms, and the number of bathrooms ranged from 1 to an outlier of 43 with an average of 2-3 bathrooms in the property. The distribution of the "acre\_lot" is highly skewed, with a mean of 10.62 acres, but a median of only 0.14, suggesting there are outliers present in the data. Similarly, there is considerable variation under "house\_size", with an average of 2,472 square feet, standard deviation of 2,326, but a maximum of 112,714. The average property price sits over \$1.27 million, with its standard deviation at over \$2.31 million, ranging from \$20,000 to \$169 million. Notably, the mean of the added variable, log-transformed price, "log\_price" is \$13.61, with a much tighter range of \$9.90 to \$18.95, highlighting how the added variable helped in normalizing the distribution of house prices, making the data more comparable to produce a more meaningful result. Lastly, from the other added variable "is\_prev\_sold", derived from the "prev\_sold\_date" variable, we found that 14.55% of properties have a recorded previous sale.






*Figure 2.1: Visual representation of the distribution of numerical features*

The histograms in Figure 2.1 each display numeric variables "bed", "bath", "acre\_lot", "house\_size", "price", and "log\_price" found in the dataset, showing that most variables are heavily skewed to the right, with a concentration of smaller values. Unsurprisingly, "log\_price" appears to be more normally distributed than other variables.





*Figure 2.2: Visual representation of the distribution of numerical features (filtered)*

After applying filters to remove extreme outliers, the updated histograms in Figure 2.2 for numeric variables "bed", "bath", "acre\_lot", "house\_size", "price", and "log\_price" now all exhibit more focused ranges, with their skewness reduced. "log\_price" remains largely unchanged. The filters applied were: removing entries with more than 15 bedrooms; removing entries with more than 15 bathrooms; removing entries with lot size larger than 2; removing entries with house size larger than 20,000 square feet; removing entries with house price higher than \$8 million.

### **b. House Features - Price**

Figure 3 is a pairplot with histograms and scatter plots. The diagonals are histograms showing distribution of bed, bath, price and log\_price. Histograms for bed, bath, and house\_size represent unimodal distribution, with clear peaks. Most properties have 3 to 6 bedrooms, with a mean of 4.5; numbers for bedrooms are rare for lower 3 and higher than 6 bedrooms. Most properties have 2 to 4 bathrooms, decreasing significantly beyond 5. House\_size histogram shows most houses are under 5,000 square feet.

Scatterplot in bed and bath reveal a positive correlation, as more bedrooms correspond to more bathrooms. Similarly larger houses will have more bedrooms and bathrooms; the 'bath and house\_size' figure shows more slope than 'bed and house\_size', indicating bathrooms influence house\_size more.

Using log\_price reduces price outliers, creates a near-normal distribution in histograms, and enhances linear correlations with bed, bath, and house\_size. As shown in the bottom row, log\_price scatter plots reveal: 1. bedrooms and log\_price show a slight positive correlation; more bedrooms will have high log\_price. However, with 10 bedrooms, log\_price varies widely (11-16). Although the bed number is the same, the log\_price does not increase within a small bound (between maximum log\_price and minimum log\_price) for a stable bedroom number, indicating other factors like bathroom number also affect price.




Figure 3: relation between bed number, bath number, and price.

2. The bed and log\_price figure shows as the number of bathrooms increases, the log\_price of properties increases. This correlation is stronger compared to bedrooms. They form a moderate positive correlation.
3. The scatterplot shows a strong positive correlation between house\_size and log\_price, with closely clustered points indicating a stable relationship less influenced by other factors. Overall, house\_size has the strongest correlation with log\_price.

## c. Geo Location - Price

### c.1. Distribution




Figure 4: Top 10 cities with highest average price & most properties.

The right side of Figure 4 shows New York City and Staten Island hold the most properties for sale than other cities', with both cities achieving over 7,000. Brooklyn and Bronx are the second tier city group containing the most properties for sale around 6,500 after New York City and Staten Island. After the fifth place, Yonkers, the number of properties listed for sale drastically decreased, in places such as Scarsdale, Jamaica, and White Plains, each holding around 1,000 properties respectively. This suggests that the top 4 cities have the highest liquidity in terms of property trading. The left side of the figure shows Old Westbury has the highest average price of properties for sale at over \$5.5, far surpassing all other cities.

## c.2. In-NYC vs Non-NYC





Figure 5: NYC House Prices vs. Non-NYC House Prices

Figure 5 shows a violin plot that compares the log-transformed house prices in New York City (NYC) and those outside of NYC. The X-axis puts the properties that are located in NYC (1 = Yes) against those that are not (0 = No), while the Y-axis displays the log-transformed prices, ranging from 10 to 16. The plot shows that the log-transformed prices of properties in NYC are generally higher than those outside of NYC. The price distribution of non-NYC houses is wider than that of NYC houses, suggesting a bigger difference between the price floor and ceiling. The median price of NYC houses is higher than the median price of non-NYC houses, as shown by the thicker part of the violin plot being in the higher price range of houses in NYC.

#### d. Previous Sale Status - Price

##### d.1. Number of Properties Sold Over Time



*Figure 6: Number of Properties Sold Over Time*

Figure 6 shows a line plot displaying the trend in the number of properties sold from the 1970s to the 2020s. Over time, the real estate market grew steadily, with an increasing number of properties being sold. However, two significant declines stand out. The first drop occurred around 2008, which aligns with the global financial crisis, a period that severely disrupted housing markets. The second decline took place around 2020, likely due to the effects of the COVID-19 pandemic. Both events highlight how external shocks can impact property sales.

##### d.2. Is\_Prev\_Sold vs Non\_Prev\_Sold






Figure 7: Violin Plot of House Prices by Previous Sales Status

In Figure 7 a violin plot was used to compare house price distributions based on whether the property had been sold before (`is_prev_sold`). The results indicate that properties with a previous sale history (`is_prev_sold = 1`) have a higher concentration of data points.

### e. Summary Correlation Matrix

Interestingly, the average prices for the two groups appear to be nearly the same. However, the upper quartile (75th percentile) of prices for properties that have not been sold before (`is_prev_sold = 0`) is slightly higher. This observation might suggest that new or unsold properties tend to achieve higher price ceilings compared to those with a prior sales history.

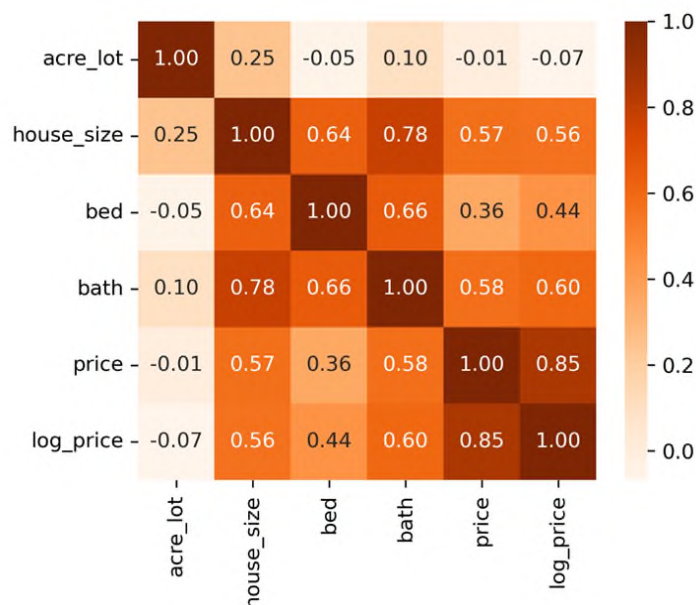


Figure 8: Correlations Between Numeric Variables

For Figure 8 we created a correlation heatmap. It provides insight into how different property features relate to house prices, particularly using `log_price` for a more normalized representation. The correlation coefficients between `log_price` and key variables are as follows: 0.44 for `bed` (number of bedrooms), 0.60 for `bath` (number of bathrooms), and 0.56 for `house_size`. These moderate positive correlations suggest that larger properties with more rooms and bathrooms generally command higher prices.

It is also worth noting that the correlation between price and `log_price` is almost identical. Moving forward, we will primarily focus on `log_price` as it scales down the variation, and tends to enhance the accuracy of further analysis.

Interestingly, the correlation between lot size (`acre_lot`) and price is very weak at **-0.07**, indicating that lot size has little impact on house prices in the dataset. This could be due to the influence of other factors, such as location, property features, or demand in the housing market. Additionally, inconsistencies are observed with `acre_lot` during further analysis, so we decided to prioritize `house_size` as a more reliable feature for modeling purposes.

Finally, we observe a strong positive correlation between `house_size` and both `bed` and `bath` (0.64 and 0.78, respectively). This aligns with expectations, as larger homes tend to accommodate more bedrooms and bathrooms.

Overall, this analysis highlights key property features that influence prices, providing a foundation for further exploration and modeling.

## **Question 2. Regression Model**

### **a. Log Transformation**

We used two variations of the dataset—one cleaned (duplicates and outliers removed) and one uncleaned (only outliers removed)—to analyze how the data sets impact model performance. This allowed us to compare the effects of both log transformation and cleaning on the accuracy and reliability of the models. We applied a log transformation to scale down the price data, reducing skewness. Alongside this, we retained an untransformed version of the price data to build a separate model, enabling a thorough comparison of the impact of transformations and data cleaning on predictive accuracy.

### **b. Variable Selection**

The selected variables encapsulate key factors influencing property prices. `price` is the target variable, while `house_size`, `bed`, and `bath` represent core features affecting functionality and appeal. `acre_lot` reflects land size, and `nyc` accounts for the unique pricing dynamics of the New York City location. This selection balances property attributes and location for a comprehensive analysis.

### **c. Model Selection**

We used the OLS linear regression model because it is a simpler and powerful model for predicting a continuous target variable like price. It assumes a linear relationship between the dependent variable (price) and the independent variables (`house_size`, `bed`, `bath`, `acre_lot`, and `nyc`). This allows us to interpret the impact of each feature on property prices, providing clear insights into how changes in features influence the price.



#### d. Model Performance

The performance of the OLS linear regression model was evaluated using key metrics. The R-squared and adjusted R-squared values of the selected model demonstrated that the model effectively captured a significant portion of the variance in property prices while accounting for the number of predictors. The p-values confirmed which key variables were statistically significant contributors to the model. Additionally, the coefficients provided clear and interpretable insights into the magnitude and direction of each variable's impact on property prices showing us which had strongest relation and which were weakest. The goodness of fit is explained in the findings below. Overall, the model performed well, delivering reliable predictions and valuable insights.

#### e. Findings

OLS Regression Results						
=====						
Dep. Variable:	price		R-squared:	0.487		
Model:	OLS		Adj. R-squared:	0.487		
Method:	Least Squares		F-statistic:	1.404e+04		
Date:	Sun, 15 Dec 2024		Prob (F-statistic):	0.00		
Time:	17:38:18		Log-Likelihood:	-1.1001e+06		
No. Observations:	74005		AIC:	2.200e+06		
Df Residuals:	73999		BIC:	2.200e+06		
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	-1.803e+05	6917.582	-26.060	0.000	-1.94e+05	-1.67e+05
nyc[T.Yes]	6.477e+05	5898.716	109.796	0.000	6.36e+05	6.59e+05
house_size	340.5055	3.465	98.263	0.000	333.714	347.297
bed	-7.48e+04	1952.312	-38.311	0.000	-7.86e+04	-7.1e+04
bath	2.201e+05	3055.031	72.052	0.000	2.14e+05	2.26e+05
acre_lot	-1.928e+05	7210.347	-26.741	0.000	-2.07e+05	-1.79e+05
=====						
Omnibus:	41683.164		Durbin-Watson:	0.132		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	548442.073		
Skew:	2.455		Prob(JB):	0.00		
Kurtosis:	15.400		Cond. No.	8.78e+03		
-----						

Figure 1 : Original Data (Excluding Outliers)

OLS Regression Results						
=====						
Dep. Variable:	log_price		R-squared:	0.509		
Model:	OLS		Adj. R-squared:	0.509		
Method:	Least Squares		F-statistic:	1.535e+04		
Date:	Sat, 14 Dec 2024		Prob (F-statistic):	0.00		
Time:	21:44:58		Log-likelihood:	-55540.		
No. Observations:	74005		AIC:	1.111e+05		
Df Residuals:	73999		BIC:	1.111e+05		
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	12.5448	0.005	2444.450	0.000	12.535	12.555
nyc[T.Yes]	0.5208	0.004	119.008	0.000	0.512	0.529
house_size	0.0002	2.57e-06	77.802	0.000	0.000	0.000
bed	0.0021	0.001	1.430	0.153	-0.001	0.005
bath	0.1699	0.002	74.953	0.000	0.165	0.174
acre_lot	-0.1896	0.005	-35.454	0.000	-0.200	-0.179
=====						
Omnibus:	5057.576		Durbin-Watson:	0.103		
Prob(Omnibus):	0.000		Jarque-Bera (JB):	13929.687		
Skew:	-0.380		Prob(JB):	0.00		
Kurtosis:	4.985		Cond. No.	8.78e+03		

Figure 2 : Original Data with Logged Prices (Excluding Outliers)

OLS Regression Results						
Dep. Variable:	log_price	R-squared:	0.494			
Model:	OLS	Adj. R-squared:	0.493			
Method:	Least Squares	F-statistic:	1545.			
Date:	Sat, 14 Dec 2024	Prob (F-statistic):	0.00			
Time:	21:49:43	Log-Likelihood:	-6693.3			
No. Observations:	7928	AIC:	1.340e+04			
Df Residuals:	7922	BIC:	1.344e+04			
Df Model:	5					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	12.4321	0.018	694.631	0.000	12.397	12.467
nyc[T.Yes]	0.5362	0.014	38.149	0.000	0.509	0.564
house_size	0.0002	8.22e-06	24.041	0.000	0.000	0.000
bed	-0.0109	0.005	-2.210	0.027	-0.021	-0.001
bath	0.2007	0.007	26.769	0.000	0.186	0.215
acre_lot	-0.2094	0.017	-12.109	0.000	-0.243	-0.175
=====						
Omnibus:	483.617	Durbin-Watson:	0.799			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1168.148			
Skew:	-0.371	Prob(JB):	2.19e-254			
Kurtosis:	4.728	Cond. No.	8.87e+03			
=====						

*Figure 3 : Cleaned Data with Logged Prices*

The analysis compares three regression models to identify which one performs best at predicting property prices. Each model approached the dataset differently, with variations in data preprocessing and transformations of the price variable, aiming to address potential issues like skewness, outliers, and duplicate records.

The first model was based on the original dataset with outliers in property prices removed but without any transformations applied to the price variable. This model achieved an R-squared value of 48.7%, meaning that 48.7% of the variation in property prices was explained by the selected predictors (*house\_size*, *bed*, *bath*, *acre\_lot*, and whether the property was in NYC). While this is a reasonable level of explanatory power, the lack of transformation may have left issues like skewed prices and heteroscedasticity unaddressed, which can limit the model's effectiveness and stability.

In the second model, the same dataset was used, but the price variable was log-transformed. This transformation helped scale down the values. As expected, the R-squared value increased to 50.9%, and the adjusted R-squared also stood at 50.9%. This adjustment reflects the proportion of variance explained by the model while accounting for the number of predictors, confirming that the improvement is meaningful. By log-transforming the price, the model captured a stronger and more linear relationship between the predictors and the target variable (price), resulting in better predictive accuracy.

In the third model we removed duplicate rows from the dataset in addition to applying the log transformation to prices. While this additional cleaning step improved the integrity of the data, the R-squared value for this model dropped slightly to 49.4%. This suggests that the removal of duplicates reduced some of the variability in the dataset that the model could use, leading to a marginally lower explanatory power. While this model still performed well, it did not surpass the second model in terms of fit or clarity.

Based on this comparison, the second model proves as the best-performing model, with an R-squared of 50.9%. This implies that the log transformation effectively addressed potential issues with the price


variable, allowing for a better fit. The adjusted R-squared being identical to the R-squared confirms that the predictors used in the model were highly relevant and contribute meaningfully to explaining the variance in property prices.

As we have chosen the second model for further analysis, a deeper observation reveals that most predictors - `house_size`, `bath`, `acre_lot`, and the NYC indicator - demonstrate strong statistical significance, with p-values consistently reported as 0.00. This means these variables have a significant impact on property prices, allowing us to confidently reject the null hypothesis that they have no effect. However, one important detail stands out: the predictor `bed` does not meet the same level of significance. Its p-value is 0.13, which is above the conventional threshold of 0.05, indicating that the number of bedrooms does not independently contribute to explaining property prices in this model.

In examining the coefficients from the chosen model (second model), the strongest predictor of property prices is the NYC indicator, with a coefficient of approximately 0.52 in the log-transformed model. This suggests that properties in the NYC area are associated with a price increase of about 52% compared to non-NYC properties, highlighting the premium of urban locations. Bathrooms also play a significant role, with a coefficient of 0.17, indicating a 17% increase in price for each additional bathroom. Meanwhile, house size contributes positively to price, though its influence is comparatively weaker, with a coefficient of 0.0002. Bedrooms and lot size have even smaller impacts, with coefficients of 0.0021 and -0.19, respectively. Interestingly, the negative coefficient for lot size suggests that larger lots might slightly reduce property prices after controlling for other factors, which could reflect market preferences for smaller, more centrally located lots in certain areas.

In conclusion, the second model, which applied a log transformation to the price variable, was selected as the most effective and reliable model for predicting property prices. It demonstrated the highest explanatory power, with an R-squared of 50.9%, and showcased a strong relationship between predictors and property prices. The findings emphasize the importance of location, amenities like bathrooms, and house size as key drivers of property value.

In conclusion, we developed and evaluated three regression models to predict property prices, each using different approaches to data preparation and transformation. After analyzing the goodness of fit, model performance, and the significance of predictors, we found that the log-transformed model using the uncleaned dataset provided the best results. This model effectively captured the relationships between property prices and key factors like location, house size, and amenities. By integrating data preprocessing with statistical analysis, we were able to build a model that offers clear insights into the drivers of property value and reliable predictions for real estate pricing.



*Figure 4: Comparison of Actual and Predicted Log-Prices*

The scatterplot compares the actual log-transformed property prices (y-axis) with the predicted values (x-axis) from our regression model, with the red dashed line representing perfect predictions ( $y = x$ ). The plot shows a clear upward trend, indicating that the model captures the overall relationship between the predictors and log prices. However, there are notable deviations from this ideal line, particularly for higher predicted values.

The scatter widens as the predicted log prices increase, suggesting the presence of heteroscedasticity. This means that the model's accuracy decreases for higher-priced properties. Points below the red line reflect underpredictions, where the actual values are higher than predicted, while points above indicate overpredictions.

While the model provides a good basis and successfully captures the general trend, the heteroscedasticity observed in the residuals suggests areas for improvement. Addressing this issue could involve refining the feature set, or exploring more advanced regression techniques to stabilize the variance and improve overall prediction accuracy,

### **Question 3. Classification Model**

#### **a.Objective:**

The goal of the project was to develop and assess different classification models like Logistic Regression, Decision Tree and Random Forest model that can predict whether a property is located in NYC or not, given that its characteristics are provided in the framework of the assignment.

As per the requirement, all geographical predictors such as "city" and "zip code" were excluded from the

analysis. The approach has been focused on structured feature selection, experimentation with different models, and thorough performance evaluation of the models. A systematic approach focusing on feature selection, model experimentation, and performance evaluation was adopted.

## **b.Methodology:**

The data cleaning process ensured that the information was relevant to the project and cleaned which was required for the comprehensive data reprocessing.

Geographical dictators like city, zip\_code, state, status, and prev\_sold\_date were removed to ensure a reliable dataset was produced for the analysis through different predicting models.

Rows with missing values in the target variable nyc were dropped and the missing values in numeric predictors - bed, bath, acre\_lot, house\_size, price were filled using their respective medians.

To determine the most suitable model for the task, three machine learning algorithms were tested:

1. Logistic Regression; 2. Decision tree; 3. Random forest

## **c. Data Preprocessing:**

### **Missing Values Handling:**

Missing value handling was one of the crucial steps involved in building the models.

In this analysis, the following steps were executed:

Missing values in nyc variables were removed for data integrity and to avoid any potential biases.

The median values for each column were used for imputing the missing values in bed, bath, acre\_lot, house\_size, and price numerical columns. Due to this approach, the central tendency is maintained while limiting the distortion from outliers.

### **Encoding Target Variable:**

Firstly the variable targets 'nyc' were converted to binary numeric for better modelling.

Feature Scaling was used to make variables more consistent and to allow the best performance of gradient-based models: Numerical predictors (bed, bath, acre\_lot, house\_size, and price) were standardized using StandardScaler.

## **d. Model Development:**

Below are the three machine learning models that were used in the analysis to predict whether a property is in NYC or not:

1. Logistic Regression; 2. Decision tree; 3. Random forest

### **Cross-Validation:**

Stratified K-Fold Cross-Validation with 5 folds was used to make a fair and balanced splits between the NYC and non-NYC properties. By using this method, the same proportion of NYC and non-NYC properties was produced in each fold, hence making more accurate results when testing these models.

### **Model Performance Evaluation Metrics:**

To assess the performance of the classification models, the following evaluation methods were used:

### **1. Confusion Matrices:**

Confusion matrices gave a detailed breakdown of the performance of the different models by showing the counts of true positives, true negatives, false positives, and false negatives. This analysis helped evaluate how well the models performed across both NYC and non-NYC property classifications.

### **2. ROC Curves:**

ROC curves were plotted to visually show the trade-off between the true positive rate (TPR) and the false positive rate (FPR). These plots showed a clear view of how well the models were classified by the two classes at different decision thresholds, thus making the task of comparing their effectiveness easier. These metrics allowed gaining valuable insight into the strengths and weaknesses of each model regarding its classification capabilities.

### **3. ROC-AUC:**

It is defined as the ability of the model to distinguish between NYC and non-NYC properties across various thresholds.

**4. Accuracy:** The overall number of correct predictions made by the model.

## **e. Results:**

The different classification models developed for predicting whether a property is in NYC showed varied performances.

The strengths and limitations of each of the models are discussed below.

### **1. Logistic Regression:**

Overall the Logistic Regression performed well, establishing itself as a good baseline model. However, some limitations in the robustness were seen. As it was sensitive to class imbalances, which affected its generalization ability.

### **2. Decision Tree:**

The Decision Tree model was highly comprehensible with its straightforward and transparent structure.

It helped in shedding light on how each feature contributed to the prediction in an understandable manner, as the relationship between different variables was recognisable.

However, it was very prone to overfitting because its high variance resulted in poor generalization between different data splits.

### **3. Random Forest:**

Random Forest proved to be the best performance model during this analysis. Being an ensemble method, it was also resistant to overfitting, with the highest accuracy and ROC-AUC score when compared to other models.

The Random Forest was the best predicting model at classifying NYC versus non-NYC properties. This fact is further supported by its confusion matrix, which has high classification



accuracy for both true positives and true negatives. The Random forest model performed the best as it handled the scaled and the balanced data well.

### Key Insights:

The results show that without the direct location data, the variable set of property size, price, and room count was sufficient to effectively classify properties in NYC.

Of all the models tested, the Random Forest was the most accurate and reliable choice for this task.


Its robustness and ability to leverage key features like house size and price make it the most suitable model for predicting the location of properties of NYC.

### Visualizations:

#### ROC Curves:

The ROC curves below illustrate the trade-off between sensitivity and specificity for each model:

##### a. Logistic Regression




*Figure 1: Logistic Regression ROC Curve*

The ROC curve for the Logistic Regression model represents a trade-off between the sensitivity (True Positive Rate) and specificity (False Positive Rate). The ROC curve is above the diagonal baseline, meaning that the model does perform better than random guessing. Area under the ROC curve is 0.78, reflecting moderate predictive performance.

The Logistic Regression model did a decent job distinguishing between NYC and non-NYC properties and an AUC of 0.78 suggests there is further work to be done to achieve even higher classification performance.

This indicates the model is reliable; however it may not be as robust as the other models.


##### b. Decision Tree



*Figure 2: Decision Tree ROC Curve*

The ROC curve for the Decision Tree model illustrates the balance between True Positive Rate(sensitivity) and False Positive Rate(specificity). The ROC curve is well above the baseline, which confirms that the model is doing better than random guessing. The AUC is 0.79, which is better compared to Logistic Regression and moderate effectiveness in distinguishing between NYC and non-NYC properties. The Decision Tree model showed better predictive performance with an AUC of 0.79. That means this model was pretty effective at classifying properties, but it is also understandable. While the curve is indicative of great performance, the model is probably prone to overfitting; this decreases its generalization capability when compared to a robust ensemble method.

### c. Random Forest



*Figure 3: Random Forest ROC Curve*

The ROC curve for the Random Forest model speaks very highly for its excellent prediction capability. The ROC curve touches the upper boundary, which indicates perfect separation by the model between the two classes (NYC versus non-NYC properties). Area under the curve equals 1.00, meaning perfect classification and no compromise between True Positive Rate and False Positive Rate. An AUC score of 1.00 emphasizes the high performance of the

Random Forest model, with no single error in the separation of NYC properties from the rest. This result underscores the reliability, strength, and suitability of the model for this kind of classification, thus making it optimal for the problem at hand.

#### d. Confusion Matrix (Random Forest):

Below is the confusion matrix of the Random Forest model:




Figure 4: Confusion Matrix Random Forest Model

The confusion matrix for the Random Forest model provides valuable insights into its predictive performance of the model.

Key results are as follows:

1. True Negatives: 12,383 properties correctly classified as not being in NYC.
2. False Positives: 48 properties incorrectly classified as not being in NYC.
3. False Negatives: 51 properties incorrectly classified as being in NYC.
4. True Positives: 4,326 properties accurately classified as being in NYC.

The Random Forest model is exhibiting very good performance, meaning that the model is correctly classifying NYC and non-NYC properties, which showcases its efficiency in handling the classification task.

The low number of misclassified properties is indicative of the robustness and reliability of the model. In general, these results confirm that the Random Forest model is the best choice for this classification task, since it presents high predictive accuracy with reliable performance.

#### e. Conclusions:

The analysis of different classification models shows that the Random Forest is the best model for classifying NYC properties. It attained the highest accuracy and ROC-AUC score, which indicates a better capability in distinguishing between NYC and non-NYC properties. The ensemble-based approach followed by Random Forest made it resistant to overfitting, hence helping it to extract meaningful patterns from the dataset. Its consistent performance across all evaluation metrics makes it the most reliable choice for this classification task.

### **1. Key Features:**

The features `house_size` and `price` showed up as significant predictors throughout all the models as per the assignments framework. These variables were very vital in differentiating NYC properties because they were strongly correlated with the value of the property.

### **2. Final Takeaway:**

Among all the predicting models, the Random Forest model was best suited for considering the problems of performance and handling diverse relationships within the dataset.

The model gives high importance to features like `house_size` and `price` to capture unique aspects of NYC properties, producing results with better accuracy in each subsequent trial.

The model can have enhanced precision and robustness for wider applications by tuning parameter adjustments and addressing class imbalances.

The results show that even in the absence of direct geographical information, property attributes carry sufficient information to classify NYC properties effectively.

## **f. Recommendations:**

To improve the models of the Random Forest, Logistic Regression and the Decision Tree the class imbalance can be dealt with by using different techniques as follows:

### **1. Hyperparameter Tuning:**

Further optimization of the Random Forest model, by tuning parameters such as the number of trees or maximum depth, may boost its performance.

### **2. Handling Imbalanced Classes:**

Techniques such as oversampling the minority class or applying weighted loss functions enhance the performance of Logistic Regression and Decision Tree models, respectively, by handling class imbalance efficiently.

### **3. Refining the model:**

The mentioned steps will help in improving the models to make those models more accurate and robust in their predictions.

## **g. Business Implications:**

The developed model provides a strong foundation for predicting the location of NYC properties with key features such as house size and price. The insights derived will help to identify the most influential factors that determine whether a property is in NYC.

These key factors will enable businesses and investors to focus their investments in properties that house these particular characteristics.

This helps businesses in making more data-driven decisions for targeting properties in NYC while optimizing their strategies.

The fact that the model can differentiate between NYC properties without the explicit use of location data demonstrates its strength and practical application in real estate decision-making.

In general, these can lead to wiser investments, better resource allocation, and increased profitability in the competitive NYC property market.

AI Declaration: AI tools were used to refine language for the report.

## Import libraries

```
##### To read and manipulate data
import pandas as pd
import numpy as np

##### To run regressions
import statsmodels.api as sm
import statsmodels.formula.api as smf
import statsmodels.graphics.api as smg

##### Various sci-kit learn functions
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score, auc
from itertools import combinations
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score, RocCurveDisplay

# Import accuracy_score
from sklearn.metrics import accuracy_score

##### For plotting
import seaborn as sns
import matplotlib.pyplot as plt
import matplotlib.colors as mcolors
import scipy.stats as stats

%pylab inline

%pylab is deprecated, use %matplotlib inline and import the required
libraries.
Populating the interactive namespace from numpy and matplotlib
```

## Color Settings

```
palette = sns.color_palette() # Default color palette
print(palette) # Prints the RGB tuples that make up this color
palette
sns.palplot(palette) # Plotting your palette!
pairpalette = sns.color_palette('Paired')
```



```
sns.palplot(pairpalette) # Seaborn color palette, with 10 colors
sns.color_palette("Oranges", as_cmap=True) # Get a CMap
```

```
[(0.12156862745098039, 0.4666666666666667, 0.7058823529411765), (1.0,
0.4980392156862745, 0.054901960784313725), (0.17254901960784313,
0.6274509803921569, 0.17254901960784313), (0.8392156862745098,
0.15294117647058825, 0.1568627450980392), (0.5803921568627451,
0.403921568627451, 0.7411764705882353), (0.5490196078431373,
0.33725490196078434, 0.29411764705882354), (0.8901960784313725,
0.4666666666666667, 0.7607843137254902), (0.4980392156862745,
0.4980392156862745, 0.4980392156862745), (0.7372549019607844,
0.7411764705882353, 0.13333333333333333), (0.09019607843137255,
0.7450980392156863, 0.8117647058823529)]
```



```
# Convert RGB tuples to hex
hex_colors = [mcolors.to_hex(color) for color in
sns.color_palette('Paired')]
print(hex_colors)

['#a6cee3', '#1f78b4', '#b2df8a', '#33a02c', '#fb9a99', '#e31a1c',
'#fdbf6f', '#ff7f00', '#cab2d6', '#6a3d9a', '#ffff99', '#b15928']
```

## Question 1: Summary Statistics

### Read Files

```
# The file path starts from the same location of this notebook
file_path = 'realtor-data-ny.csv'
df = pd.read_csv(file_path)
```

### 1. Data Cleaning

First, look at the data

```
# Create the 'prev_sold' variable as a explanation for entries with
their 'prev_sold_date' variable missing. This shows whether a house
had been previously sold.
df['is_prev_sold'] = df['prev_sold_date'].notnull().astype(int)

# Create the log transformed 'price' variable, 'log_price' for better
relative comparison, since entieres are spread across different
geographic locations, with exceptional anomalies displayed in New York
City.
df['log_price'] = np.log(df['price'])
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 84040 entries, 0 to 84039
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   status                84040 non-null  object
1   bed                   84040 non-null  int64
2   bath                  83946 non-null  float64
3   acre_lot              84040 non-null  float64
4   city                  84038 non-null  object
5   state                 84040 non-null  object
6   zip_code              84036 non-null  float64
7   house_size            84040 non-null  int64
8   prev_sold_date        56605 non-null  object
9   price                 84040 non-null  int64
10  nyc                   84040 non-null  object
11  is_prev_sold          84040 non-null  int32
12  log_price             84040 non-null  float64
dtypes: float64(4), int32(1), int64(3), object(5)
memory usage: 8.0+ MB
```

```
df.head()
```

	status	bed	bath	acre_lot	city	state	zip_code
house_size \							
0 for_sale	3	1.0	0.37	Accord	New York	12404.0	
960							
1 for_sale	3	2.0	0.38	Accord	New York	12404.0	
1936							
2 for_sale	2	1.0	0.41	Accord	New York	12404.0	
832							
3 for_sale	3	1.0	5.50	Accord	New York	12404.0	
1900							
4 for_sale	3	3.0	6.50	Accord	New York	12404.0	
4000							
	prev_sold_date	price	nyc	is_prev_sold	log_price		

0	21/03/2022	249900	No	1	12.428816
1	06/01/1989	319000	No	1	12.672946
2	10/09/2015	169500	No	1	12.040608
3	NaN	695000	No	0	13.451667
4	30/07/2021	250000	No	1	12.429216

```
df.tail()
```

	status	bed	bath	acre_lot	city	state	zip_code
house_size \							
84035	for_sale	3	3.0	0.65	Yulan	New York	12792.0
1480							
84036	for_sale	4	1.0	1.64	Yulan	New York	12792.0
1692							
84037	for_sale	4	1.0	1.64	Yulan	New York	12792.0
1692							
84038	for_sale	5	2.0	0.13	NaN	New York	NaN
1925							
84039	for_sale	2	1.0	110.00	NaN	New York	12523.0
1177							

	prev_sold_date	price	nyc	is_prev_sold	log_price
84035	23/05/2006	425000	No	1	12.959844
84036	20/01/2012	188500	No	1	12.146853
84037	20/01/2012	188500	No	1	12.146853
84038	NaN	710000	No	0	13.473020
84039	NaN	495000	No	0	13.112313

```
df.describe()
```

	bed	bath	acre_lot	zip_code
house_size \				
count	84040.000000	83946.000000	84040.000000	84036.000000
84040.000000				
mean	3.930200	2.980690	10.624121	10983.073409
2472.279938				
std	2.062923	1.756449	849.058032	688.870753
2326.090138				
min	1.000000	1.000000	0.000000	6390.000000
122.000000				
25%	3.000000	2.000000	0.060000	10514.000000
1370.000000				
50%	4.000000	3.000000	0.140000	10916.000000
2000.000000				
75%	5.000000	4.000000	0.570000	11233.000000
2880.000000				
max	42.000000	43.000000	100000.000000	14534.000000
112714.000000				

	price	is_prev_sold	log_price
--	-------	--------------	-----------

count	8.404000e+04	84040.000000	84040.000000
mean	1.274604e+06	0.673548	13.612734
std	2.312462e+06	0.468917	0.836738
min	2.000000e+04	0.000000	9.903488
25%	5.280000e+05	0.000000	13.176852
50%	7.545000e+05	1.000000	13.533811
75%	1.220000e+06	1.000000	14.014361
max	1.690000e+08	1.000000	18.945409

```
df.isnull().sum()
```

```
status      0
bed          0
bath        94
acre_lot     0
city         2
state        0
zip_code     4
house_size   0
prev_sold_date  27435
price        0
nyc          0
is_prev_sold  0
log_price    0
dtype: int64
```

For the records that are missing 'Prev\_sold\_data', we decided that it just means there's no record that the property's been sold. It's still valid data, so we kept them.

## 1.1 distribution of categorical variables

```
# List of variables to analyze
variables = ['status', 'city', 'state', 'prev_sold_date',
            'is_prev_sold', 'nyc']
# Loop through variables and display value counts
for col in variables:
    print(f"Distribution of values for {col}:")
    print(df[col].value_counts())
    print("\n")
```

Distribution of values for status:

```
status
for_sale    84040
Name: count, dtype: int64
```

Distribution of values for city:

```
city
New York City    8338
Staten Island    7073
```

```
Brooklyn      6535
Bronx         6424
Yonkers       2321
...
Godeffroy     1
Grahamsville  1
Northport     1
Hewlett Harbor 1
Staatsburg    1
Name: count, Length: 457, dtype: int64
```

```
Distribution of values for state:
state
New York      84040
Name: count, dtype: int64
```

```
Distribution of values for prev_sold_date:
prev_sold_date
04/11/2003     134
05/11/2021     104
21/06/2017      81
03/10/2013      77
06/09/2005      73
...
20/03/2008      1
10/07/1998      1
12/07/2017      1
11/07/2006      1
21/03/2022      1
Name: count, Length: 4267, dtype: int64
```

```
Distribution of values for is_prev_sold:
is_prev_sold
1      56605
0      27435
Name: count, dtype: int64
```

```
Distribution of values for nyc:
nyc
No      61936
Yes     22104
Name: count, dtype: int64
```

## 1.2 Distribution of continuous variables

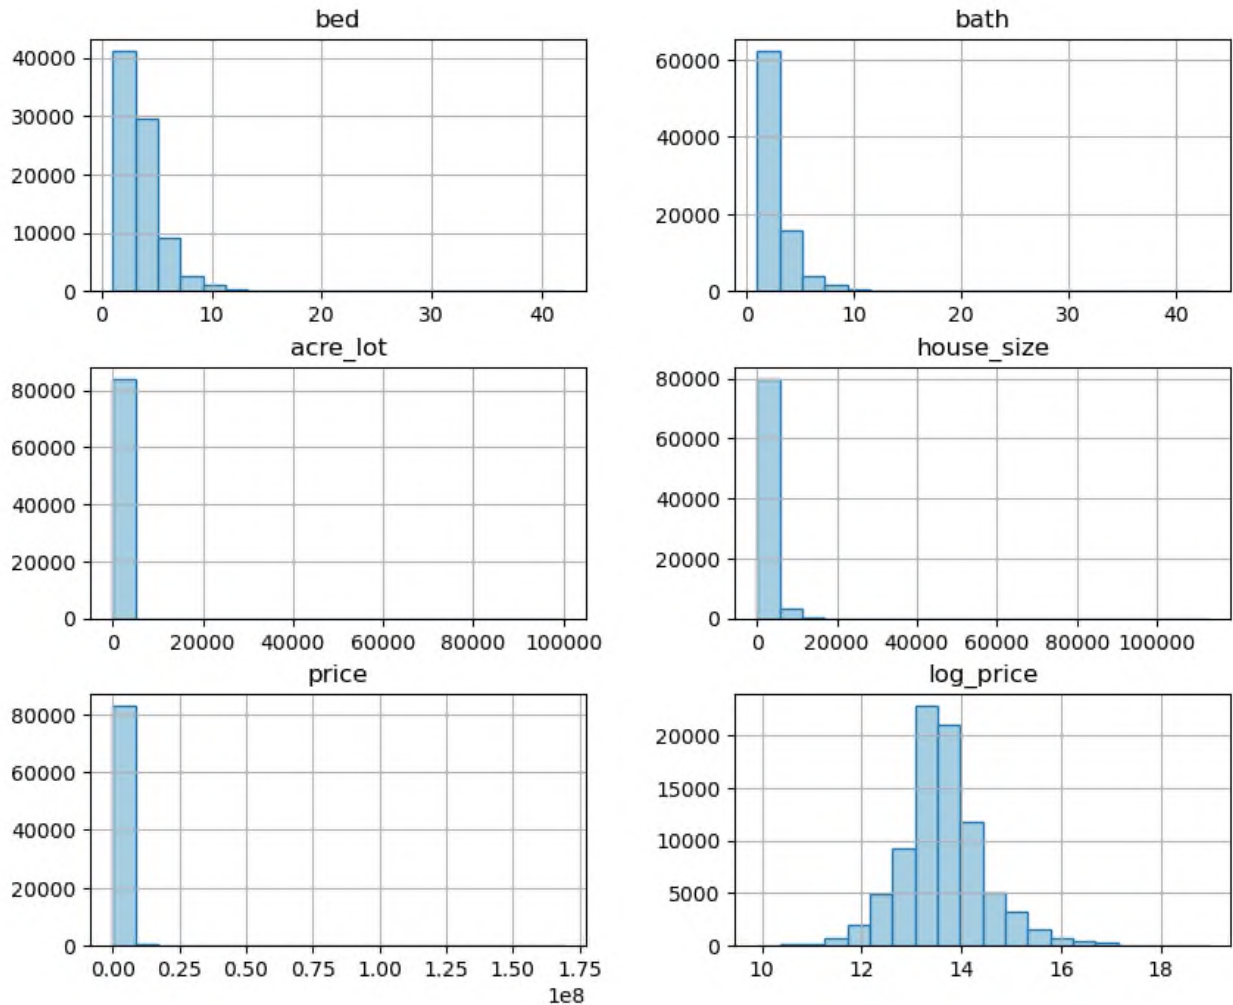
```
# List of columns for which to plot histograms
columns_to_plot = ['bed', 'bath', 'acre_lot', 'house_size', 'price',
'log_price']
print(df[columns_to_plot].describe())
# Plot histograms for the specified columns
df[columns_to_plot].hist(figsize=(10,8), bins=20,
color=pairpalette[0],edgecolor=pairpalette[1])
```

	bed	bath	acre_lot	house_size
price \				
count	84040.000000	83946.000000	84040.000000	84040.000000
8.404000e+04				
mean	3.930200	2.980690	10.624121	2472.279938
1.274604e+06				
std	2.062923	1.756449	849.058032	2326.090138
2.312462e+06				
min	1.000000	1.000000	0.000000	122.000000
2.000000e+04				
25%	3.000000	2.000000	0.060000	1370.000000
5.280000e+05				
50%	4.000000	3.000000	0.140000	2000.000000
7.545000e+05				
75%	5.000000	4.000000	0.570000	2880.000000
1.220000e+06				
max	42.000000	43.000000	100000.000000	112714.000000
1.690000e+08				

	log_price
count	84040.000000
mean	13.612734
std	0.836738
min	9.903488
25%	13.176852
50%	13.533811
75%	14.014361
max	18.945409

```
array([[<Axes: title={'center': 'bed'}>,
      <Axes: title={'center': 'bath'}>],
      [<Axes: title={'center': 'acre_lot'}>,
      <Axes: title={'center': 'house_size'}>],
      [<Axes: title={'center': 'price'}>,
      <Axes: title={'center': 'log_price'}>]], dtype=object)
```





## Remove Outliers

*# Apply all filtering conditions at once*

```
cdf = df[
    (df['bed'] <= 15) &
    (df['bath'] <= 15) &
    (df['acre_lot'] <= 2) &
    (df['house_size'] <= 20000) &
    (df['price'] <= 8000000)
]
```

*# Check the resulting data*

```
print(cdf.describe())
```

	bed	bath	acre_lot	zip_code
house_size \				
count	74005.000000	74005.000000	74005.000000	74001.000000
74005.000000				
mean	3.882724	2.829836	0.294322	10968.352806

2221.820107				
std	1.849067	1.418526	0.392463	654.365575
1283.951529				
min	1.000000	1.000000	0.000000	6390.000000
122.000000				
25%	3.000000	2.000000	0.060000	10512.000000
1360.000000				
50%	4.000000	3.000000	0.110000	10923.000000
1950.000000				
75%	5.000000	3.000000	0.360000	11233.000000
2713.000000				
max	14.000000	15.000000	2.000000	14534.000000
15000.000000				


	price	is_prev_sold	log_price
count	7.400500e+04	74005.000000	74005.000000
mean	1.033275e+06	0.682805	13.567891
std	9.643704e+05	0.465387	0.731436
min	2.000000e+04	0.000000	9.903488
25%	5.390000e+05	0.000000	13.197471
50%	7.500000e+05	1.000000	13.527828
75%	1.178000e+06	1.000000	13.979329
max	7.999000e+06	1.000000	15.894827

*# Plot cleaned data*

```
cdf[columns_to_plot].hist(figsize=(10,8), color=pairpalette[0],
bins=20, edgecolor=pairpalette[1])
```

```
array([[<Axes: title={'center': 'bed'}>,
        <Axes: title={'center': 'bath'}>],
       [<Axes: title={'center': 'acre_lot'}>,
        <Axes: title={'center': 'house_size'}>],
       [<Axes: title={'center': 'price'}>,
        <Axes: title={'center': 'log_price'}>]], dtype=object)
```





### 3. Geo Location vs. Price

#### 3.1 Top 10 Cities with highest average price


```
# Group by 'city' and calculate the mean 'price'
avg_price_by_city = cdf.groupby('city')['price'].mean()

# Sort the results by average price, in descending order
sorted_avg_price_by_city =
avg_price_by_city.sort_values(ascending=False)

# Select the top 10 cities
top_10_avg_price_by_city = sorted_avg_price_by_city.head(10)
```

```
top_10_avg_price_by_city.plot(kind='bar', figsize=(5, 4),
color=pairpalette[0], title='Top 10 Cities with Highest Average
Price', xlabel='City', ylabel='Average Price')

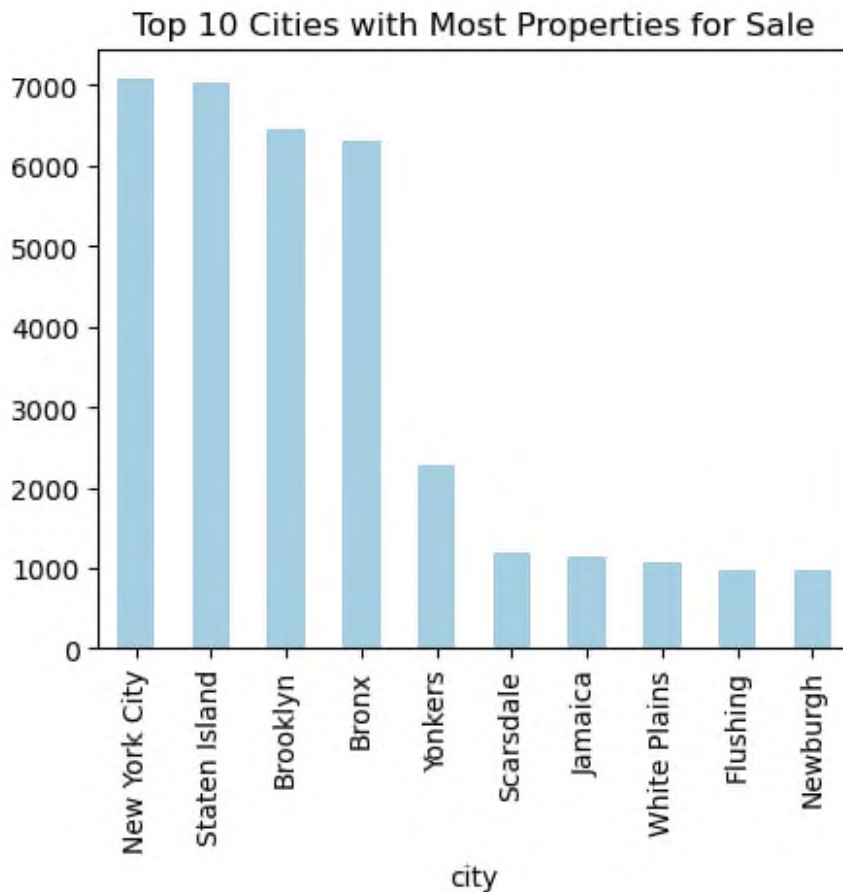
<Axes: title={'center': 'Top 10 Cities with Highest Average Price'},
xlabel='City', ylabel='Average Price'>
```



### 3.2 Top 10 Cities with Most Properties for Sale

```
cdf['city'].value_counts().head(10).plot(kind='bar', figsize=(5,4),
color=pairpalette[0], title='Top 10 Cities with Most Properties for
Sale')

<Axes: title={'center': 'Top 10 Cities with Most Properties for
Sale'}, xlabel='city'>
```



### 3.3 NYC vs. Non-NYC

```
print(cdf['nyc'].value_counts())
```


```
nyc
No      53294
Yes     20711
Name: count, dtype: int64
```

```
df['nyc'].value_counts().plot(kind='pie', figsize=(5, 4),
    autopct='%1.0f%%', startangle=90, color=[pairpalette[1],
    pairpalette[7]], title='Number of Properties: NYC vs Non-NYC')
```


```
<Axes: title={'center': 'Number of Properties: NYC vs Non-NYC'},
    ylabel='count'>
```



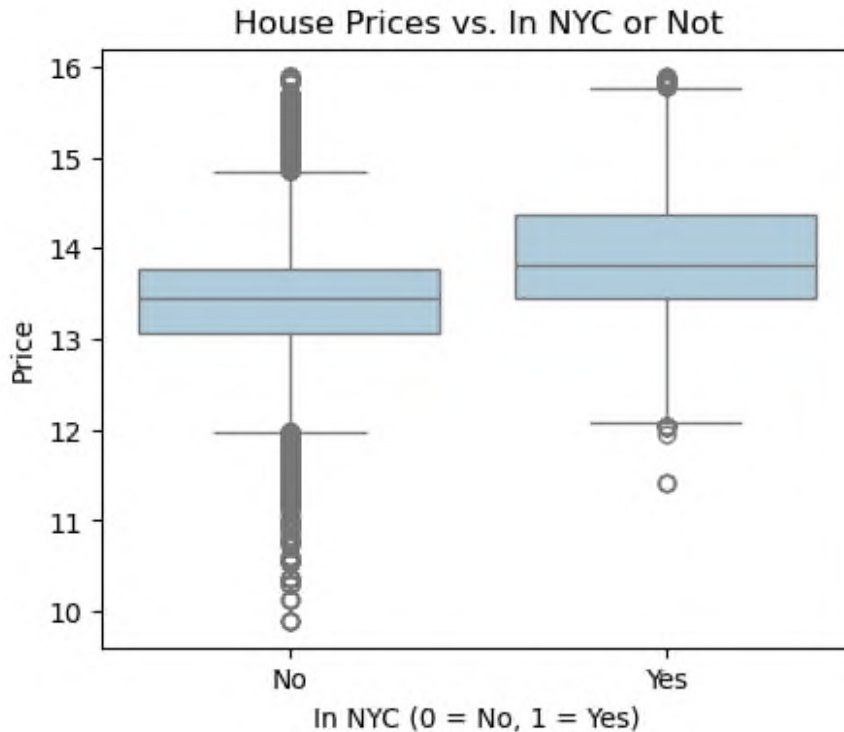
Number of Properties: NYC vs Non-NYC



```
# Create a violin plot
plt.figure(figsize=(5,4))
sns.violinplot(x='nyc', y='log_price', data=cdf,
inner='quartile',edgecolor=pairpalette[1],color=pairpalette[0])
plt.title('NYC House Prices vs. Non-NYC House Prices')
plt.xlabel('In NYC (0 = No, 1 = Yes)')
plt.ylabel('Price(Log Scale)')
Text(0, 0.5, 'Price(Log Scale)')
```



```
# Create a box plot
plt.figure(figsize=(5,4))
sns.boxplot(x='nyc', y='log_price', data=cdf,color=pairpalette[0])
plt.title('House Prices vs. In NYC or Not')
plt.xlabel('In NYC (0 = No, 1 = Yes)')
plt.ylabel('Price')
plt.show()
```



## 4. Status - Price

### 4.1 Properties Sold Over Time


```
df['prev_sold_date'] = pd.to_datetime(df['prev_sold_date'],
errors='coerce')
df['prev_sold_date'].dt.year.value_counts().sort_index().plot(kind='line',
figsize=(8,4), title='Number of Properties Sold Over Time')

# Annotations: Housing Market Crash (2008) and COVID-19 (2020)
plt.axvline(x=2008, color='red', linestyle='--', label='2008 Housing
Crash', alpha=0.4)
plt.axvline(x=2020, color='blue', linestyle='--', label='COVID-19
Impact', alpha=0.4)
plt.legend()

# Adding text annotations
plt.text(2008, 200, '08 Crash', color='red', ha='center', va='center',
fontsize=10, alpha=0.8)
plt.text(2020, 300, 'COVID-19', color='blue', ha='center',
va='center', fontsize=10, alpha=0.8)
plt.show()
```

C:\Users\gsh\AppData\Local\Temp\ipykernel\_12056\464379496.py:1:  
UserWarning: Parsing dates in %d/%m/%Y format when dayfirst=False (the  
default) was specified. Pass `dayfirst=True` or specify a format to  
silence this warning.

```
df['prev_sold_date'] = pd.to_datetime(df['prev_sold_date'],
errors='coerce')
```



## 4.2 Is\_pre\_sold vs on\_pre\_sold

```
print(cdf['is_prev_sold'].value_counts())
```

```
is_prev_sold
```

```
1    50531
```

```
0    23474
```

```
Name: count, dtype: int64
```


```
df['is_prev_sold'].value_counts().plot(kind='bar', figsize=(4, 3),
color=[pairpalette[1], pairpalette[7]])
```

```
plt.title('Average House Prices by Previous Sale Status')
```


```
plt.xlabel('Is Previously Sold (0 = No, 1 = Yes)')
```

```
plt.ylabel('Average Price')
```


```
plt.show()
```



```
# Create a violin plot
plt.figure(figsize=(5,4))
sns.violinplot(x='is_prev_sold', y='log_price', data=cdf,
inner='quartile',edgecolor=pairpalette[1],color=pairpalette[0])
plt.title('House Prices vs. Previous Sale Status')
plt.xlabel('Is Previously Sold (0 = No, 1 = Yes)')
plt.ylabel('Price')
Text(0, 0.5, 'Price')
```



```
# Create a box plot
plt.figure(figsize=(5,4))
sns.boxplot(x='is_prev_sold', y='log_price',
data=cdf,color=pairpalette[0])
plt.title('House Prices vs. Previous Sale Status')
plt.xlabel('Is Previously Sold (0 = No, 1 = Yes)')
plt.ylabel('Price')
plt.show()
```

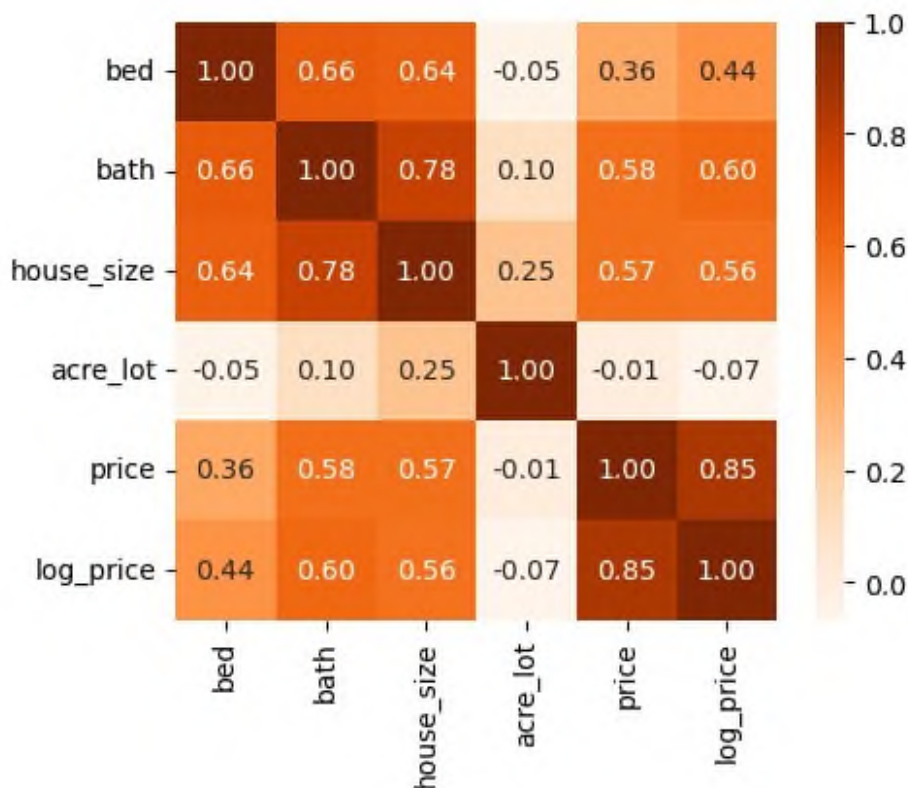


## 5. Summary & Others

### 5.1 Summary Correlation Heatmap

```
plt.figure(figsize=(5,4))
sns.heatmap(cdf[['bed', 'bath', 'house_size', 'acre_lot', 'price', 'log_price']].corr(), annot=True, cmap='Oranges', fmt='.2f')
```

<Axes: >



## Question 2: Regression Model

```
model = smf.ols(formula= 'price ~ house_size+bed+bath+acre_lot+nyc',
data=cdf).fit()
print(model.summary())
```

### OLS Regression Results

```
=====
=====
Dep. Variable:                price    R-squared:
0.487
Model:                        OLS      Adj. R-squared:
0.487
Method:                       Least Squares    F-statistic:
1.404e+04
Date:                         Mon, 16 Dec 2024    Prob (F-statistic):
0.00
Time:                         11:30:08    Log-Likelihood:
1.1001e+06
No. Observations:             74005    AIC:
2.200e+06
Df Residuals:                 73999    BIC:
2.200e+06
Df Model:                     5
```



Covariance Type: nonrobust

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept    -1.803e+05    6917.582     -26.060      0.000    -1.94e+05    -
1.67e+05
nyc[T.Yes]   6.477e+05    5898.716     109.796      0.000     6.36e+05
6.59e+05
house_size    340.5055         3.465      98.263      0.000      333.714
347.297
bed          -7.48e+04    1952.312     -38.311      0.000    -7.86e+04
-7.1e+04
bath          2.201e+05    3055.031      72.052      0.000     2.14e+05
2.26e+05
acre_lot     -1.928e+05    7210.347     -26.741      0.000    -2.07e+05    -
1.79e+05
=====
=====
Omnibus:                41683.164    Durbin-Watson:
0.132
Prob(Omnibus):           0.000    Jarque-Bera (JB):
548442.073
Skew:                    2.455    Prob(JB):
0.00
Kurtosis:                15.400    Cond. No.
8.78e+03
=====
=====
```

#### Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 8.78e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
model = smf.ols(formula= 'log_price ~
house_size+bed+bath+acre_lot+nyc', data=cdf).fit()
print(model.summary())
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          log_price    R-squared:
```

0.509  
Model: OLS Adj. R-squared:  
0.509  
Method: Least Squares F-statistic:  
1.535e+04  
Date: Mon, 16 Dec 2024 Prob (F-statistic):  
0.00  
Time: 11:30:08 Log-Likelihood:  
-55540.  
No. Observations: 74005 AIC:  
1.111e+05  
Df Residuals: 73999 BIC:  
1.111e+05  
Df Model: 5

Covariance Type: nonrobust

```
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept      12.5448        0.005    2444.450      0.000      12.535
12.555
nyc[T.Yes]      0.5208        0.004    119.008      0.000      0.512
0.529
house_size      0.0002      2.57e-06     77.802      0.000      0.000
0.000
bed              0.0021        0.001      1.430      0.153     -0.001
0.005
bath             0.1699        0.002     74.953      0.000      0.165
0.174
acre_lot        -0.1896        0.005    -35.454      0.000     -0.200
-0.179
=====
=====
```

Omnibus: 5057.576 Durbin-Watson:  
0.103  
Prob(Omnibus): 0.000 Jarque-Bera (JB):  
13929.687  
Skew: -0.380 Prob(JB):  
0.00  
Kurtosis: 4.985 Cond. No.  
8.78e+03  
=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is

correctly specified.

[2] The condition number is large,  $8.78e+03$ . This might indicate that there are strong multicollinearity or other numerical problems.

```
# Plot residuals vs fitted values
```

```
fitted_values = model.fittedvalues # Predicted values
```

```
residuals = model.resid # Residuals (actual - predicted)
```

```
plt.figure(figsize=(6,4))
```

```
sns.scatterplot(x=fitted_values, y=residuals)
```


```
plt.axhline(0, color='red', linestyle='--') # Reference line at y=0
```

```
plt.title('Residuals vs Fitted Values')
```

```
plt.xlabel('Fitted Values (Predicted log_price)')
```

```
plt.ylabel('Residuals')
```

```
plt.show()
```




```
# Q-Q plot
```

```
plt.figure(figsize=(6,4))
```


```
stats.probplot(residuals, dist="norm", plot=plt)
```

```
plt.title('Q-Q Plot of Residuals')
```


```
plt.show()
```



```
plt.figure(figsize=(6,4))
sns.histplot(residuals, kde=True, bins=30)
plt.title('Histogram of Residuals')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.show()
```



```
# Actual vs Predicted
plt.figure(figsize=(6,4))
sns.scatterplot(x=model.fittedvalues, y=cdf['log_price'])
plt.plot(model.fittedvalues, model.fittedvalues, color='red',
linestyle='--') # Ideal line (y = x)
plt.title('Actual vs Predicted Values')
plt.xlabel('Predicted log_price')
plt.ylabel('Actual log_price')
plt.show()
```



```
drop_dup_cdf = cdf.drop_duplicates()
model = smf.ols(formula= 'log_price ~
house_size+bed+bath+acre_lot+nyc', data=drop_dup_cdf).fit()
print(model.summary())
```

#### OLS Regression Results

```
=====
=====
Dep. Variable:          log_price    R-squared:
0.494
Model:                  OLS         Adj. R-squared:
0.493
Method:                 Least Squares    F-statistic:
1545.
Date:                   Mon, 16 Dec 2024    Prob (F-statistic):
0.00
Time:                   11:30:10    Log-Likelihood:
-6693.3
No. Observations:      7928    AIC:
1.340e+04
Df Residuals:          7922    BIC:
1.344e+04
Df Model:              5
Covariance Type:      nonrobust
```

```

=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
Intercept      12.4321      0.018     694.631      0.000      12.397
12.467
nyc[T.Yes]      0.5362      0.014      38.149      0.000      0.509
0.564
house_size      0.0002     8.22e-06      24.041      0.000      0.000
0.000
bed             -0.0109      0.005      -2.210      0.027     -0.021
-0.001
bath            0.2007      0.007      26.769      0.000      0.186
0.215
acre_lot       -0.2094      0.017     -12.109      0.000     -0.243
-0.175
=====
=====
Omnibus:                483.617   Durbin-Watson:
0.799
Prob(Omnibus):          0.000   Jarque-Bera (JB):
1168.148
Skew:                   -0.371   Prob(JB):
2.19e-254
Kurtosis:               4.728   Cond. No.
8.87e+03
=====
=====

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is
correctly specified.
[2] The condition number is large, 8.87e+03. This might indicate that
there are
strong multicollinearity or other numerical problems.

```

### Question 3: Classification Model

```

# Drop rows with missing 'nyc' values
df = df.dropna(subset=['nyc'])

# Save the cleaned DataFrame to a new CSV file
df.head()

```

	status	bed	bath	acre_lot	city	state	zip_code
house_size \							
0 for_sale	3	1.0	0.37	Accord	New York	12404.0	
960							

1	for_sale	3	2.0	0.38	Accord	New York	12404.0
1936							
2	for_sale	2	1.0	0.41	Accord	New York	12404.0
832							
3	for_sale	3	1.0	5.50	Accord	New York	12404.0
1900							
4	for_sale	3	3.0	6.50	Accord	New York	12404.0
4000							

	prev_sold_date	price	nyc	is_prev_sold	log_price
0	2022-03-21	249900	No	1	12.428816
1	1989-01-06	319000	No	1	12.672946
2	2015-09-10	169500	No	1	12.040608
3	NaT	695000	No	0	13.451667
4	2021-07-30	250000	No	1	12.429216

df.tail()

	status	bed	bath	acre_lot	city	state	zip_code
house_size \							
84035	for_sale	3	3.0	0.65	Yulan	New York	12792.0
1480							
84036	for_sale	4	1.0	1.64	Yulan	New York	12792.0
1692							
84037	for_sale	4	1.0	1.64	Yulan	New York	12792.0
1692							
84038	for_sale	5	2.0	0.13	NaN	New York	NaN
1925							
84039	for_sale	2	1.0	110.00	NaN	New York	12523.0
1177							

	prev_sold_date	price	nyc	is_prev_sold	log_price
84035	2006-05-23	425000	No	1	12.959844
84036	2012-01-20	188500	No	1	12.146853
84037	2012-01-20	188500	No	1	12.146853
84038	NaT	710000	No	0	13.473020
84039	NaT	495000	No	0	13.112313

```
# Handle missing values in numeric columns by filling with the median
numeric_columns = ['bed', 'bath', 'acre_lot', 'house_size', 'price']
df[numeric_columns] = df[numeric_columns].apply(lambda x:
x.fillna(x.median()))
```

```
# Encode the target variable 'nyc' (Yes -> 1, No -> 0)
le_nyc = LabelEncoder()
df['nyc'] = le_nyc.fit_transform(df['nyc'])
df.head()
```

	status	bed	bath	acre_lot	city	state	zip_code
house_size \							
0	for_sale	3	1.0	0.37	Accord	New York	12404.0



```

960
1  for_sale      3    2.0      0.38  Accord  New York    12404.0
1936
2  for_sale      2    1.0      0.41  Accord  New York    12404.0
832
3  for_sale      3    1.0      5.50  Accord  New York    12404.0
1900
4  for_sale      3    3.0      6.50  Accord  New York    12404.0
4000

```

```

      prev_sold_date  price  nyc  is_prev_sold  log_price
0      2022-03-21  249900    0             1  12.428816
1      1989-01-06  319000    0             1  12.672946
2      2015-09-10  169500    0             1  12.040608
3              NaT  695000    0             0  13.451667
4      2021-07-30  250000    0             1  12.429216

```

```

# Define the feature set
features = ['bed', 'bath', 'acre_lot', 'house_size', 'price']

# Scale the features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df[features])
y = df['nyc']

# Initialize classifiers
log_reg = LogisticRegression(random_state=42)
decision_tree = DecisionTreeClassifier(random_state=42)
rf_model = RandomForestClassifier(random_state=42)

# Define Stratified K-Fold cross-validation (5 folds)
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Perform Stratified K-Fold cross-validation for Logistic Regression
log_reg_scores = cross_val_score(log_reg, X_scaled, y, cv=skf,
scoring='accuracy')
print("Logistic Regression Stratified K-Fold Accuracy Scores:",
log_reg_scores)
print(f"Logistic Regression Mean Accuracy:
{np.mean(log_reg_scores):.4f}")

# Perform Stratified K-Fold cross-validation for Decision Tree
decision_tree_scores = cross_val_score(decision_tree, X_scaled, y,
cv=skf, scoring='accuracy')
print("\nDecision Tree Stratified K-Fold Accuracy Scores:",
decision_tree_scores)
print(f"Decision Tree Mean Accuracy:
{np.mean(decision_tree_scores):.4f}")

# Perform Stratified K-Fold cross-validation for Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

```

```

rf_scores = cross_val_score(rf_model, X_scaled, y, cv=skf,
scoring='accuracy')

print("\nRandom Forest Stratified K-Fold Accuracy Scores:", rf_scores)
print(f"Random Forest Mean Accuracy: {np.mean(rf_scores):.4f}")

Logistic Regression Stratified K-Fold Accuracy Scores: [0.78807711
0.78873156 0.78825559 0.78891004 0.78908853]
Logistic Regression Mean Accuracy: 0.7886

Decision Tree Stratified K-Fold Accuracy Scores: [0.99280105 0.993396
0.99303903 0.99262256 0.99113517]
Decision Tree Mean Accuracy: 0.9926

Random Forest Stratified K-Fold Accuracy Scores: [0.99369348
0.99375297 0.99393146 0.99327701 0.99440743]
Random Forest Mean Accuracy: 0.9938

# Placeholder function for evaluating feature combinations
def evaluate_feature_combinations(model, X, y, features):
    best_combo = None
    best_score = 0
    for r in range(1, len(features) + 1):
        for combo in combinations(features, r):
            X_combo = X[list(combo)]
            scores = cross_val_score(model, X_combo, y, cv=5,
scoring='roc_auc')
            mean_score = np.mean(scores)
            if mean_score > best_score:
                best_score = mean_score
                best_combo = combo
    best_model = model.fit(X[list(best_combo)], y)
    return best_combo, best_score, best_model

# Convert the scaled data back to a DataFrame for easy feature
selection
X_scaled_df = pd.DataFrame(X_scaled, columns=features)

# Split the data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X_scaled_df, y,
test_size=0.2, random_state=42)

# Train Random Forest model
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print(f"Random Forest Test Accuracy: {rf_accuracy:.2f}")

Random Forest Test Accuracy: 0.99

```

```

# Placeholder for best_rf_combo and best_rf_score
best_rf_combo = features
best_rf_score = rf_accuracy
best_rf_model = rf_model

# Evaluate Logistic Regression
print("Evaluating Logistic Regression...")
best_log_reg_combo, best_log_reg_score, best_log_reg_model =
evaluate_feature_combinations(
    log_reg, X_scaled_df, y, features
)

Evaluating Logistic Regression...

# Evaluate Decision Tree
print("\nEvaluating Decision Tree...")
best_tree_combo, best_tree_score, best_tree_model =
evaluate_feature_combinations(
    decision_tree, X_scaled_df, y, features
)

Evaluating Decision Tree...

# Evaluate Random Forest model
rf_predictions = rf_model.predict(X_test)
rf_accuracy = accuracy_score(y_test, rf_predictions)
print(f"Random Forest Accuracy: {rf_accuracy:.2f}")

Random Forest Accuracy: 0.99

# Display the best results
print("\nBest Logistic Regression Feature Combination:",
best_log_reg_combo, "with ROC AUC:", best_log_reg_score)
print("Best Decision Tree Feature Combination:", best_tree_combo,
"with ROC AUC:", best_tree_score)
print("Best Random Forest Feature Combination:", best_rf_combo, "with
ROC AUC:", best_rf_score)

Best Logistic Regression Feature Combination: ('bed', 'acre_lot',
'house_size', 'price') with ROC AUC: 0.7677090297237096
Best Decision Tree Feature Combination: ('acre_lot',) with ROC AUC:
0.6814030705289911
Best Random Forest Feature Combination: ['bed', 'bath', 'acre_lot',
'house_size', 'price'] with ROC AUC: 0.9941099476439791

def plot_roc_curve(model, X, y, model_name, feature_combo):

    # Ensure feature_combo is a list, not a tuple
    feature_combo = list(feature_combo)

```

```

# Predict probabilities for the positive class
y_pred_prob = model.predict_proba(X[feature_combo])[:, 1]


# Calculate the ROC curve
fpr, tpr, thresholds = roc_curve(y, y_pred_prob)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.figure(figsize=(5, 4))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve
(area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title(f'ROC Curve - {model_name}')
plt.legend(loc="lower right")
plt.show()


# Plot ROC curves for both models
plot_roc_curve(best_log_reg_model, X_scaled_df, y, "Logistic
Regression", best_log_reg_combo)
plot_roc_curve(best_tree_model, X_scaled_df, y, "Decision Tree",
best_tree_combo)
plot_roc_curve(best_rf_model, X_scaled_df, y, "Random Forest",
best_rf_combo)

```


<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>



<Figure size 640x480 with 0 Axes>



```
# Confusion Matrix for Random Forest
rf_cm = confusion_matrix(y_test, rf_predictions)
plt.figure(figsize=(4, 4))
sns.heatmap(rf_cm, annot=True, fmt='d', cmap='Blues')
plt.title("Random Forest Confusion Matrix")
plt.ylabel("Actual")
plt.xlabel("Predicted")
plt.show()
```

